

FAI LAB 4

Beyond Classical Search

Paolo Morettin

2024-25


Code available

You can find the **code** of the lab sessions on GitHub:

<https://github.com/paolomorettin/FAI-code>


On non-determinism

- In a computer: (seed number +) pseudo-random algorithms
- We need to agree on a **weighted sampling** procedure



On non-determinism

- In a computer: (seed number +) pseudo-random algorithms
- We need to agree on a **weighted sampling** procedure




- The idea: a list of pre-determined choices $\in [0, 1]$ is provided

$choices = [0.5, 0.1, 0.9]$
Selected: *B*

On non-determinism

- In a computer: (seed number +) pseudo-random algorithms
- We need to agree on a **weighted sampling** procedure




- The idea: a list of pre-determined choices $\in [0, 1]$ is provided

$choices = [0.5, 0.1, 0.9]$
Selected: **A**

On non-determinism

- In a computer: (seed number +) pseudo-random algorithms
- We need to agree on a **weighted sampling** procedure




- The idea: a list of pre-determined choices $\in [0, 1]$ is provided

$choices = [0.5, 0.1, \textcolor{red}{0.9}]$
Selected: D

On non-determinism

- In a computer: (seed number +) pseudo-random algorithms
- We need to agree on a **weighted sampling** procedure



- The idea: a list of pre-determined choices $\in [0, 1]$ is provided

$choices = [0.5, 0.1, \textcolor{red}{0.9}]$
Selected: D

- Then we cycle

On non-determinism

```
def select_choice(options, choices):
    '''Simulated stochastic process that deterministically pick an
    option given a pre-determined list of choices. Choices are
    cycled through. Options are weighted.

    ...
    choice = choices.pop(0)
    choices.append(choice) # cycling

    w_sum = sum(w for _, w in options)
    p = 0
    for opt, w in options:
        p += w/w_sum
        if choice <= p:
            return opt, p
```

Local search: hill climbing


- Goal: find s^* maximizing an objective function $f(\cdot)$
- One rule: **do not look down**

candidates $H(curr) = \{s \in Neigh(curr) \mid f(s) > f(curr)\}$

- Different variants:
 - **steepest**: $next = \text{argmax}_{n \in H(curr)} f(n)$
 - **stochastic (unweighted)**: $next \sim \mathcal{U}(H(curr))$
 - **stochastic**: $next \sim p(n) \propto f(H(curr))$
- Multiple (parallel) restarts of the above

Non-convex optimization on a 2D grid

- **Goal:** maximize a randomly generated function $f(x, y)$



- $\text{Actions}((x, y))$ results in the following order of next states:

$[(x - 1, y - 1), (x, y - 1), (x + 1, y - 1),$
 $(x - 1, y), (x + 1, y),$
 $(x - 1, y + 1), (x, y + 1), (x + 1, y + 1)]$

6	7	8
4		5
1	2	3

Local search: genetic algorithms

- Well-known family of **evolutionary algorithms**

Local search: genetic algorithms

- Well-known family of **evolutionary algorithms**
- The objective function f is called *fitness*

Local search: genetic algorithms

- Well-known family of **evolutionary algorithms**
- The objective function f is called *fitness*
- Each solution is encoded as a string over a finite alphabet

Local search: genetic algorithms

- Well-known family of **evolutionary algorithms**
- The objective function f is called *fitness*
- Each solution is encoded as a string over a finite alphabet
- At each iteration, a **population** is considered:

Local search: genetic algorithms

- Well-known family of **evolutionary algorithms**
- The objective function f is called *fitness*
- Each solution is encoded as a string over a finite alphabet
- At each iteration, a **population** is considered:
 - Individuals are selected for reproduction (with p proportional to f)

Local search: genetic algorithms

- Well-known family of **evolutionary algorithms**
- The objective function f is called *fitness*
- Each solution is encoded as a string over a finite alphabet
- At each iteration, a **population** is considered:
 - Individuals are selected for reproduction (with p proportional to f)
 - The offspring is the result of the combination of the parents

Local search: genetic algorithms

- Well-known family of **evolutionary algorithms**
- The objective function f is called *fitness*
- Each solution is encoded as a string over a finite alphabet
- At each iteration, a **population** is considered:
 - Individuals are selected for reproduction (with p proportional to f)
 - The offspring is the result of the combination of the parents
 - Mutations can happen with small probability

Local search: genetic algorithms

- Well-known family of **evolutionary algorithms**
- The objective function f is called *fitness*
- Each solution is encoded as a string over a finite alphabet
- At each iteration, a **population** is considered:
 - Individuals are selected for reproduction (with p proportional to f)
 - The offspring is the result of the combination of the parents
 - Mutations can happen with small probability
 - Elitism: retain best scoring individuals from the previous generation, cull the weak f monotonically increases

GA example: the quest for the Master Sandwitch

GA example: the quest for the Master Sandwitch

- The **DNA** of a sandwitch:

<i>main</i>	<i>side</i>	<i>sauce</i>	<i>bread</i>
H am	L ettuce	M ayo	B un
S alami	T omato	Y oghurt	W rap
F alafel	O nions	G arlic	P ita
K ebab	B ell peppers		

(e.g. *HLMP* = a ham / lettuce / mayo / pita sandwich)

GA example: the quest for the Master Sandwitch

- The **DNA** of a sandwitch:

<i>main</i>	<i>side</i>	<i>sauce</i>	<i>bread</i>
H am	L ettuce	M ayo	B un
S alami	T omato	Y oghurt	W rap
F alafel	O nions	G arlic	P ita
K ebab	B ell peppers		

(e.g. $HLMP$ = a ham / lettuce / mayo / pita sandwich)

- Let's assume a simple fitness function:

$$f(x) = \begin{cases} 0 & \text{if } x_0 = H \\ 1 & \text{if } x_0 = S \\ 2 & \text{if } x_0 = F \\ 3 & \text{if } x_0 = K \end{cases} + \begin{cases} 0 & \text{if } x_1 = L \\ 1 & \text{if } x_1 = T \\ 2 & \text{if } x_1 = O \\ 3 & \text{if } x_1 = B \end{cases} + \begin{cases} 0 & \text{if } x_2 = M \\ 1 & \text{if } x_2 = Y \\ 2 & \text{if } x_2 = G \end{cases} + \begin{cases} 0 & \text{if } x_3 = B \\ 1 & \text{if } x_3 = W \\ 2 & \text{if } x_3 = P \end{cases}$$

(It is an **oversimplification** anyway, no bacon considered)

Online search

- **Interleaving** *computation and action*


- **Interleaving** *computation* and *action*
- Useful for **dynamic** task environments, but also **nondeterministic**

- **Interleaving** *computation* and *action*
- Useful for **dynamic** task environments, but also **nondeterministic**
- Agent knows the following:
 - $\text{Actions}(s)$, what can be done
 - $\text{IsGoal}(s)$, of course
 - $s' = \text{Result}(s, a)$ not known in advance
 - $\text{cost}(s, a, s')$ known only after executing a

- **Interleaving** computation and action
- Useful for **dynamic** task environments, but also **nondeterministic**
- Agent knows the following:
 - $\text{Actions}(s)$, what can be done
 - $\text{IsGoal}(s)$, of course
 - $s' = \text{Result}(s, a)$ not known in advance
 - $\text{cost}(s, a, s')$ known only after executing a
- We focus on **safely explorable** state spaces
a goal state is reachable from any reachable state (no dead-ends)

Maze problems

- **Goal:** reach the green room
- Order of $Actions(s) = [N, E, S, W]$



- undirected graph → reversible actions
- $cost(s, a, s') = cost(s', a', s)$

Online DFS

```
function ONLINE-DFS-AGENT(problem, s') returns an action
    s, a, the previous state and action, initially null
    persistent: result, a table mapping (s, a) to s', initially empty
    untried, a table mapping s to a list of untried actions
    unbacktracked, a table mapping s to a list of states never backtracked to

    if problem.IS-GOAL(s') then return stop
    if s' is a new state (not in untried) then untried[s']  $\leftarrow$  problem.ACTIONS(s')
    if s is not null then
        result[s, a]  $\leftarrow$  s'
        add s to the front of unbacktracked[s']
    if untried[s'] is empty then
        if unbacktracked[s'] is empty then return stop
        else a  $\leftarrow$  an action b such that result[s', b] = POP(unbacktracked[s'])
    else a  $\leftarrow$  POP(untried[s'])
    s  $\leftarrow$  s'
    return a
```

Learning Real Time A*

function LRTA*-AGENT(*problem*, s' , h) **returns** an action
 s , a , the previous state and action, initially null
persistent: *result*, a table mapping (s, a) to s' , initially empty
 H , a table mapping s to a cost estimate, initially empty

if IS-GOAL(s') **then return** *stop*
if s' is a new state (not in H) **then** $H[s'] \leftarrow h(s')$
if s is not null **then**
 $result[s, a] \leftarrow s'$
 $H[s] \leftarrow \min_{b \in \text{ACTIONS}(s)} \text{LRTA}^{\text{-COST}}(s, b, result[s, b], H)$
 $a \leftarrow \operatorname{argmin}_{b \in \text{ACTIONS}(s)} \text{LRTA}^{\text{-COST}}(\text{problem}, s', b, result[s', b], H)$
 $s \leftarrow s'$
return a

function LRTA*-COST(*problem*, s, a, s', H) **returns** a cost estimate
if s' is undefined **then return** $h(s)$
else return *problem.ACTION-COST*(s, a, s') + $H[s']$