

COURSE "AUTOMATED PLANNING: THEORY AND PRACTICE"

CHAPTER 07: GENERAL SEARCH STRATEGIES

Teacher: **Marco Roveri** - marco.roveri@unitn.it

M.S. Course: Artificial Intelligence Systems (LM)

A.A.: 2025-2026

Where: DISI, University of Trento

URL: <https://shorturl.at/A81hf>



Last updated: Sunday 12th October, 2025

TERMS OF USE AND COPYRIGHT

USE

This material (including video recording) is intended solely for students of the University of Trento registered to the relevant course for the Academic Year 2025-2026.

SELF-STORAGE

Self-storage is permitted only for the students involved in the relevant courses of the University of Trento and only as long as they are registered students. Upon the completion of the studies or their abandonment, the material has to be deleted from all storage systems of the student.

COPYRIGHT

The copyright of all the material is held by the authors. Copying, editing, translation, storage, processing or forwarding of content in databases or other electronic media and systems without written consent of the copyright holders is forbidden. The selling of (parts) of this material is forbidden. Presentation of the material to students not involved in the course is forbidden. The unauthorised reproduction or distribution of individual content or the entire material is not permitted and is punishable by law.

The material (text, figures) in these slides is authored by Jonas Kvarnström and Marco Roveri.

IMPORTANT DISTINCTION

OPTIMIZING

- Optimal plan generation:
 - There is a **quality measure** for plans
 - (Minimal number of actions)
 - Minimal **sum of action costs**
 - ...
- We **must** find an optimal plan!
 - Suboptimal plans (0.5% more expensive): **irrelevant!**

Guaranteeing optimality is sometimes **useful**, always **expensive**!

SATISFICING

- Satisficing (satisfy/suffice) in general:
 - *"Searching until an acceptability threshold is met"*
 - Motivation: High-quality non-optimal solutions are also useful
 - Can often be found in reasonable time
- Satisficing in **planning** (typically):
 - No well-defined threshold: **Any form of non-optimal planning**
 - *Try to find strategies and heuristics that seem reasonably quick and give reasonable results in our tests*

Investigate many **different points** on the efficiency/quality spectrum!

INFORMED VS UNINFORMED SEARCH

UNINFORMED SEARCH

- No domain-specific knowledge
- Can only take into account **search space structure** and **cost so far**
 - $g(n)$ = cost of reaching node n from a starting point

INFORMED SEARCH

- Take additional information into account, such as heuristics!

Applicable to all search spaces we have seen so far

May work *better* in some of them...

DIJKSTRA'S ALGORITHM

- Matches the forward search "template"
 - Use a "*simple*" strategy to select and remove a node n from open
 - Select a node n with minimal $g(n)$: Cost of reaching n from initial node
- Efficient graph search algorithm: $O(|E| + |V| \log(|V|))$
 - $|E|$ = number of edges (transitions), $|V|$ = number of nodes (states)

```


function SEARCH(problem)
    initial-node ← MAKE-INITIAL-NODE(problem)                                → [2]
    open ← {initial-node}
    while (open ≠ ∅) do
        node ← SEARCH-STRATEGY-REMOVE-FROM(open)                            → [6]
        if IS-SOLUTION(node) then                                         → [4]
            return EXTRACT-PLAN-FROM(node)                                     → [5]
        end if
        ...
    end while
    ...
end function

```

Typical Implementation
Priority Queue


DIJKSTRA'S ALGORITHM: EXPLORATION ORDER

- Explore nodes in increasing/decreasing order of cost!




DIJKSTRA'S ALGORITHM: TOWER OF HANOI

- Running Dijkstra, assuming all ToH actions are equally expensive




DIJKSTRA'S ALGORITHM: BLOCKS WORLD

- Running Dijkstra, assuming all BW actions are equally expensive



DIJKSTRA'S ALGORITHM: EXAMPLE

- A small instance



DIJKSTRA'S ALGORITHM: EXAMPLE

- A typical implementation: 8706 created states, 2692 visited/expanded
- BW 400
 - Standard formulation: $s^{n^2+3n+1} = 2^{161201} > 10^{48526}$ states
 - But we do not have to visit every one ... fewer reachable states!
- BW 400 - blocks initially on the table, goal is a 400-block tower
 - Given state space search with uniform action costs (same cost for all actions), Dijkstra will always consider all plans that stack less than 400 blocks!
 - Stacking 1 block: = plans, 400×399 plans, ...
 - Stacking 2 blocks: $> 400 \times 399 \times 399 \times 398$ plans, ...
 - Will visit more than 1.63×10^{1735}

Dijkstra is efficient in terms of search space size: $O(|E| + |V| \log(|V|))$

The search space is exponential in the size of the input description...

FAST COMPUTERS, MANY CORES

- But computers are getting **very fast**!
 - Suppose we can check 10^{20} states per second
 - > 10 billion states per clock cycle for today's computers, each state involving complex operations
 - Then it will only take $10^{1735}/10^{20} = 10^{1715}$ seconds..
- But we have **multiple cores**!
 - The universe has at most 10^{87} particles, including electrons, ...
 - Let's suppose every one is a CPU core
 - \Rightarrow only 10^{1628} seconds $> 10^{1620}$ years!
 - The universe is around 10^{10} years old!



IMPRACTICAL ALGORITHMS

- Dijkstra's algorithm is **completely impractical** here
 - Visits all nodes with $cost < cost(optimal\ solution)$
- If we don't guarantee optimality: **Depth first search?**
 - Could be faster, by pure luck...
 - but normally finds **very** inefficient plans

The state space is fine, but we need some *guidance*

BEST FIRST SEARCH: INTUITION

```

function SEARCH(problem)
    initial-node  $\leftarrow$  MAKE-INITIAL-NODE(problem)
    open  $\leftarrow$  {initial-node}
    while (open  $\neq \emptyset$ ) do
        node  $\leftarrow$  SEARCH-STRATEGY-REMOVE-FROM(open)
        if IS-SOLUTION(node) then
            return EXTRACT-PLAN-FROM(node)
        end if
        for each newnode  $\in$  SUCCESSORS(node) do
            open  $\leftarrow$  open  $\cup$  {newnode}
        end for
    end while
    return Failure
end function

```

- Keep track of a set of open nodes
- Use an heuristic function $h(\text{node})$ to select the open node that seems "best"
- As opposed to depth-first, breadth-first, ... which only consider tree structure!
- As opposed to Dijkstra's algorithm etc.,.. which consider cost so far, and having no idea where to go next!
- As opposed to hill climbing and others that "throw away nodes instead of keeping all nodes in open!"

GREEDY BEST FIRST SEARCH: INTUITION

```
function SEARCH(problem)
    initial-node ← MAKE-INITIAL-NODE(problem)
    open ← {initial-node}
    while (open ≠ ∅) do
        node ← SEARCH-STRATEGY-REMOVE-FROM(open)
        if IS-SOLUTION(node) then
            return EXTRACT-PLAN-FROM(node)
        end if
        for each newnode ∈ SUCCESSORS(node) do
            open ← open ∪ {newnode}
        end for
    end while
    return Failure
end function
```

- Choose an open node **minimizing** $h(n)$

- Ignore the cost $g(n)$ of reaching the node
- Try to minimize the (apparent) amount of search left to do

A*

- Optimal Plan Generation often uses A*
 - A* focuses **entirely** in **optimality**
 - Expands from the initial node, systematically checking all possibilities
 - No point in trying to find a "reasonable" plan before finding the optimal one!
 - Requires **admissible** heuristics to guarantee optimality: $\forall n. h(n) \leq h^*(n)$
 - $h^*(n)$ cost of optimal plan from n
 - Reason: heuristic used for **pruning** (skipping some search nodes and all descendants)
- How admissibility helps?
 - Let 12 be the cost of optimal solution
 - Another node n with $g(n) = 10$ and $h(n) = 5$
 - $h(n)$ admissible, never overestimates, so any solution from here would cost at least $10+5=15$
 - **No need to investigate successors of this node!**
 - If $h(n)$ does not underestimate, it **does not help**!
 - Could find solutions of cost 10 as descendants of node $n \implies$ must keep searching!

A* STRATEGY

- Pick nodes from open in order of increasing $f(n) = g(n) + h(n)$
 - $g(n)$ actual cost
 - $h(n)$ heuristic
- Works like a priority queue

$$11 = 10 + 1$$

$$12 = 10 + 2$$

$$12 = 12 + 0$$

$$12 = 11 + 1$$

$$13 = 11 + 2$$

Pop - not a solution

Pop - not a solution

Pop - solution

Ignore the rest: g is known, h underestimates so solution found by expanding these nodes will cost $\geq g + h$ (and we have one of cost $\leq g + h$)

- If an heuristic never underestimates costs:
 - Let 12 be the cost of a solution
 - Another node, n : $g(n) = 10$, and $h(n) = 5$
 - $h(n)$ never underestimates, so any solution found from here on would cost at most 15
 - Does not help! Could find solutions of cost 10 as descendant of node n , must keep searching!

A*: DIJKSTRA's VS A* – ESSENTIAL DIFFERENCE

DIJKSTRA

- Selects from open a node n with minimal $g(n)$
 - Cost of reaching n from initial node


Uninformed - blind -

- Example:
 - Hand-coded heuristic function
 - Can move diagonally $\Rightarrow h(n) = \max(\text{abs}(n.x-g.x), \text{abs}(n.y-g.y))$
– Chebyshev distance
 - Related to Manhattan Distance = $\text{abs}(n.x - g.x) + \text{abs}(n.y - g.y)$

A*


- Selects from open a node with minimal $g(n) + h(n)$
 - + underestimate cost of reaching a goal from n

Informed



A*

- Given an admissible heuristic h , A* is **optimal** in two ways
 - Guarantee an **optimal** plan is extracted
 - Expands the **minimum number of nodes** required to *guarantee optimality* with the given heuristic!
- Still may expand many "unproductive" nodes in the maze example
 - The heuristic is **not perfectly informative**
 - Does not take **obstacles** into account
- If we knew actual remaining cost $h^*(n)$:
 - Expand optimal path to the goal!



VARIATIONS OF A*

- Weighted A*
 - Use $f(n) = g(n) + w \cdot h(n)$
 - Weight $w > 1$ place greater emphasis on being close to the goal! I.e., you *believe* to be close to the goal!
 - \Rightarrow At most w times more expensive!
- Repeated Weighted A*
 - Consider an ordered set of weights, and try to repeatedly solve problem using one weight from the set!
 - **for** $w \in \{5.0, 3.0, 2.0, 1.0\}$ **do**
 - solve problem with Weighted A* using w
 - Rationale
 - Each pass is "much" faster than the next
 - Try to *approach* optimality while still being able to *return a plan quickly* if necessary!
 - Why not a single weight? \Rightarrow Can't predict how much time any given weight will require!
- More variants are discussed in the path planning robotic course!

WITH OPEN LIST

```
function SEARCH(problem)
    initial-node ← MAKE-INITIAL-NODE(problem)
    open ← {initial-node}
    while (open ≠ ∅) do
        node ← SEARCH-STRATEGY-REMOVE-FROM(open)
        if IS-SOLUTION(node) then
            return EXTRACT-PLAN-FROM(node)
        end if
        ...
    end while
    ...
end function
```

- With an Open List, we have no "current position" during the search!
- We choose from all open nodes, not from the nearest one!

WITHOUT OPEN LIST

```
function DEPTH-FIRST-SEARCH(problem)
    initial-node  $\leftarrow$  MAKE-INITIAL-NODE(problem)
    return DEPTH-FIRST-SEARCH-REC(initial-node)
end function
```


```
function DEPTH-FIRST-SEARCH-REC(node)
    if IS-SOLUTION(node) then
        return EXTRACT-PLAN-FROM(node)
    end if
    for each newnode  $\in$  SUCCESSORS(node) do
        solution  $\leftarrow$  DEPTH-FIRST-SEARCH-REC(newnode)
        if solution  $\neq$  null then
            return solution
        end if
    end for
    return null
end function
```

- Depth First Search can use open list or recursive search!
 - We can **only** look at the successors of *current node*
 - No possibility to postponing a node until later
 - Introduces **backtracking**: going back from *where you are*
 - Not such concept exists with open list!

STEEPEST ASCENT HILL CLIMBING

- Greedy local search algorithm for optimization problems
 - (i) Start in some current location


2D EXAMPLE



<http://www.willmcginnis.com/2012/05/12/272/>

```
function STEEPESTASCENTHILLCLIMBING(problem)
  n ← initial-node
  ...
  ...
```


STATE SPACE EXAMPLE



STEEPEST ASCENT HILL CLIMBING (CONT.)

- (ii) Find the local neighborhood, with nodes that can be reached in one step


2D EXAMPLE



<http://www.willmcginnis.com/2012/05/12/272/>

```
function STEEPESTASCENTHILLCLIMBING(problem)
    n ← initial-node
    while True do
        if n is a solution then return n
        expand children of n
```

STATE SPACE EXAMPLE



STEEPEST ASCENT HILL CLIMBING (CONT.)


- (iii) Try to **improve** using **local optimal** choice:
 - Choose the successor/neighbor that is *best in this step*
 - \Rightarrow Don't care about the *future*

```

function STEEPESTASCENTHILLCLIMBING(problem)
  n  $\leftarrow$  initial-node
  while True do
    if n is a solution then return n
    expand children of n
    calculate  $h$  for children
    if some child decreases  $h(n)$  then
      n  $\leftarrow$  a child minimizing  $h(n)$ 
    else ???
    ...
  
```

- Search nodes have no **absolute** quality
 - They are *solutions* or useless *non-solutions*
- But we can *estimate* the quality using heuristics (leading towards the goal)!

STEEPEST ASCENT HILL CLIMBING: EXAMPLE



STEEPEST ASCENT HILL CLIMBING (CONT.)

```
function GREEDYBESTFIRSTSEARCH(problem)
    n ← initial-node
    open ←  $\emptyset$ 
    while True do
        if  $n$  is a solution then return  $n$ 
        expand children of  $n$ 
        calculate  $h$  for children
        add children to open
         $n \leftarrow$  a node in open minimizing  $h(n)$ 
```


```
function STEEPESTASCENTHILLCLIMBING(problem)
    n ← initial-node
    while True do
        if  $n$  is a solution then return  $n$ 
        expand children of  $n$ 
        calculate  $h$  for children
        if some child decreases  $h(n)$  then
             $n \leftarrow$  a child minimizing  $h(n)$ 
        else stop
    ...
    Chose best among childrens
```

\rightarrow Local optimum

Be stubborn: Only consider childrens of this node, don't keep track of open nodes to return to!


LOCAL OPTIMA

- (iv) When there is **nothing strictly better** nearby: Stop!
 - Standard Hill Climbing used for *optimization*
 - Any point is a *solution*: we search for a *good* one!
 - Might find a *local optimum*: the top of a hill!



LOCAL OPTIMA (CONT.)

- Classical planning \Rightarrow *absolute goals*
 - Even if we can't decrease $h(n)$, we can simply *stop*!



LOCAL OPTIMA (CONT.)


- Standard solution to local optima:
 - Randomly choose another node
 - Continue searching from there
 - Hope you find a global optimum eventually!
- In planning:
 - Must choose a node that you have actually created during expansion...

```

function STEEPESTASCENTHILLCLIMBING(problem)
  n ← initial-node
  while True do
    if n is a solution then return n
    expand children of n
    calculate  $h$  for children
    if some child decreases  $h(n)$  then
      n ← a child minimizing  $h(n)$ 
    else
      n ← some random state
  
```



HILL CLIMBING WITH h_{add} : PLATEAUS




- No successor **improves** the heuristic value: some are equal!
 - We have a **plateau**



- Jump to a random node *immediately*?
 - No! The heuristic is not so accurate – may be some child *is closer* to the goal even though $h(n)$ is not lower!
 - \Rightarrow keep exploring: allow some consecutive **moves across plateaus**!

HILL CLIMBING WITH h_{add} : LOCAL OPTIMA



- If we continue, all successors have **higher** heuristic values!
- We have a **local** optimum... *Impasse* = optimum or plateaus \Rightarrow Some *impasses* allowed!

IMPASSES AND RESTARTS

- What if there are **many** impasses?
 - May be we are in the wrong part of the search space after all....
 - \Rightarrow Select another *promising* expanded node where search continues...

HSP 1: HEURISTIC SEARCH PLANNER

- HSP 1.x: h_{add} heuristic + hill climbing + modifications

```
function STEEPESTASCENTHILLCLIMBING(problem)
```

```
    impasses  $\leftarrow$  0
```

```
    unexpanded  $\leftarrow$   $\emptyset$ 
```

```
    current  $\leftarrow$  initial-node
```

```
    while (not yet reached the goal) do
```

```
        children  $\leftarrow$  EXPAND(current)
```

```
        if (children =  $\emptyset$ ) then
```

```
            current  $\leftarrow$  POP(unexpanded)
```

```
        else
```

```
            bestChild  $\leftarrow$  BEST(children)
```

```
            add other childrens to unexpanded in order of  $h(n)$ 
```

```
            if ( $h(bestChild) \geq h(current)$ ) then
```

```
                impasses++
```

```
                if (impasses = threshold) then
```

```
                    current  $\leftarrow$  POP(unexpanded)
```

```
                    impasses  $\leftarrow$  0
```

```
                else
```

```
                    current  $\leftarrow$  bestChild
```

```
            else
```

```
                current  $\leftarrow$  bestChild
```

→ Apply all applicable actions

→ Dead end \implies restart!

→ Child with the lowest heuristic value

→ Keep for restarts!

→ Essentially HC, but not all steps have to move "up"

→ Too many downhill/plateau moves \implies escape!

→ Restart from another node!


Simple structure, but highly competitive at its introduction!

ENFORCED HILL CLIMBING

- FastForward (FF) [1] uses **enforced** hill climbing – approximately

- $s \leftarrow$ init-state
- repeat**

expand breadth-first until a better state s' is found
until a goal state is found




Wait longer to decide which branch to take! \Rightarrow Do not restart – keep going!

PROPERTIES OF ENFORCED HILL-CLIMBING

- Is Enforced Hill-Climbing **complete**?

- No!



REFERENCES I

- [1] FF. The Fast Forward Planner. <https://fai.cs.uni-saarland.de/hoffmann/ff.html>, 2001. 34
- [2] Hector Geffner and Blai Bonet. *A Concise Introduction to Models and Methods for Automated Planning*. Synthesis Lectures on Artificial Intelligence and Machine Learning. Morgan & Claypool Publishers, 2013. ISBN 9781608459698. doi: 10.2200/S00513ED1V01Y201306AIM022. URL <https://doi.org/10.2200/S00513ED1V01Y201306AIM022>.
- [3] Malik Ghallab, Dana S. Nau, and Paolo Traverso. *Automated planning - theory and practice*. Elsevier, 2004. ISBN 978-1-55860-856-6.
- [4] Malik Ghallab, Dana S. Nau, and Paolo Traverso. *Automated Planning and Acting*. Cambridge University Press, 2016. ISBN 978-1-107-03727-4. URL <http://www.cambridge.org/de/academic/subjects/computer-science/artificial-intelligence-and-natural-language-processing/automated-planning-and-acting?format=HB>.
- [5] Patrik Haslum, Nir Lipovetzky, Daniele Magazzeni, and Christian Muise. *An Introduction to the Planning Domain Definition Language*. Synthesis Lectures on Artificial Intelligence and Machine Learning. Morgan & Claypool Publishers, 2019. doi: 10.2200/S00900ED2V01Y201902AIM042. URL <https://doi.org/10.2200/S00900ED2V01Y201902AIM042>.