

COURSE "AUTOMATED PLANNING: THEORY AND PRACTICE"

CHAPTER 03: PLANNING AS SEARCH

Teacher: **Marco Roveri** - marco.roveri@unitn.it

M.S. Course: Artificial Intelligence Systems (LM)

A.A.: 2025-2026

Where: DISI, University of Trento

URL: <https://shorturl.at/A81hf>



Last updated: Wednesday 24th September, 2025

TERMS OF USE AND COPYRIGHT

USE

This material (including video recording) is intended solely for students of the University of Trento registered to the relevant course for the Academic Year 2025-2026.

SELF-STORAGE

Self-storage is permitted only for the students involved in the relevant courses of the University of Trento and only as long as they are registered students. Upon the completion of the studies or their abandonment, the material has to be deleted from all storage systems of the student.

COPYRIGHT

The copyright of all the material is held by the authors. Copying, editing, translation, storage, processing or forwarding of content in databases or other electronic media and systems without written consent of the copyright holders is forbidden. The selling of (parts) of this material is forbidden. Presentation of the material to students not involved in the course is forbidden. The unauthorised reproduction or distribution of individual content or the entire material is not permitted and is punishable by law.

The material (text, figures) in these slides is authored by Jonas Kvarnström and Marco Roveri.

HOW TO GENERATE A PLAN?

- One way of defining planning:

Using knowledge about the world, including possible actions and their results, to decide what to do and when in order to achieve an objective, before you actually start doing it

⇒ How?

IS PLANNING A STRAIGHT-FORWARD
PROCESS OF ADDING ACTIONS IN THE
RIGHT ORDER?

```
while (exists unvisited position) do
    pos ← nearest unvisited position
    plan += flyto(pos)
    plan += aim()
    plan += take-picture()
end while
```



USUALLY, CONDITIONS ARE TOO
COMPLEX; BETTER TO
TEST-ALTERNATIVES - SEARCH

```
Generate some starting point
while (not complete) do
    try some alternatives
    create modified alternatives
    throw away some alternatives
end while
```

PLANNING AS SEARCH

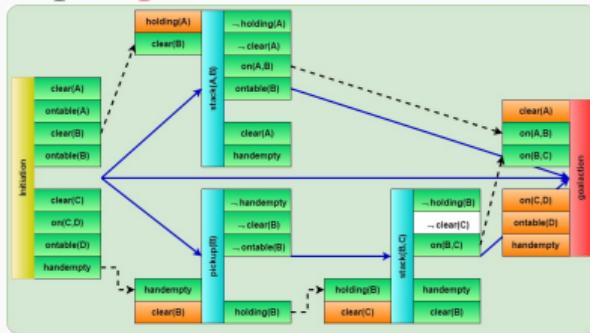
- To generate plans using search, we need:

1) A NODE STRUCTURE DEFINING WHAT INFORMATION IS IN A NODE

State space search: A node is a **state**

$$s = \{fact_1, fact_2, fact_3, \dots\}$$

Partial Order Causal Link (POCL): A node is a complex **plan structure**

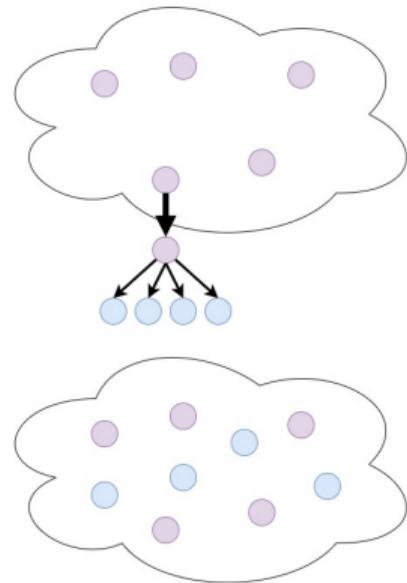


2) A WAY OF CREATING AN INITIAL NODE FROM A PROBLEM INSTANCE

- Search spaces are generally too large to be represented completely
- Need to **start somewhere** and then expand the space incrementally

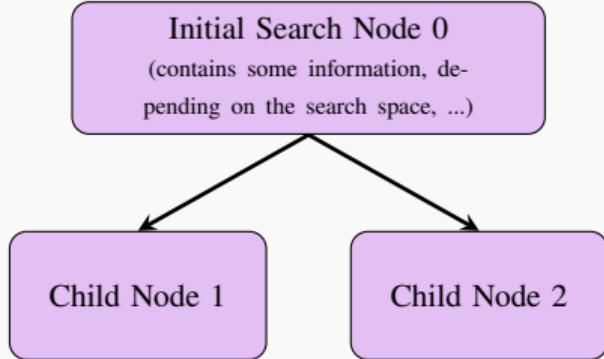
PLANNING AS SEARCH (CONT.)

- General way of **formalizing** search algorithms:
 - There are some "open" nodes (at first the **initial node**), that we
 - Know how to reach
 - Haven't explored yet
 - Pick/Remove** one of them
 - Using some *strategy* to peek "good" nodes!
 - Find **neighbor** nodes that can be created in a "single" step
 - Put **created** nodes in the set of "open" nodes
 - Repeat until a node **corresponds to a solution**



PLANNING AS SEARCH (CONT.)

SEARCH SPACE
CLASSICAL PLANNING
FINITE NUMBER OF SEARCH NODES

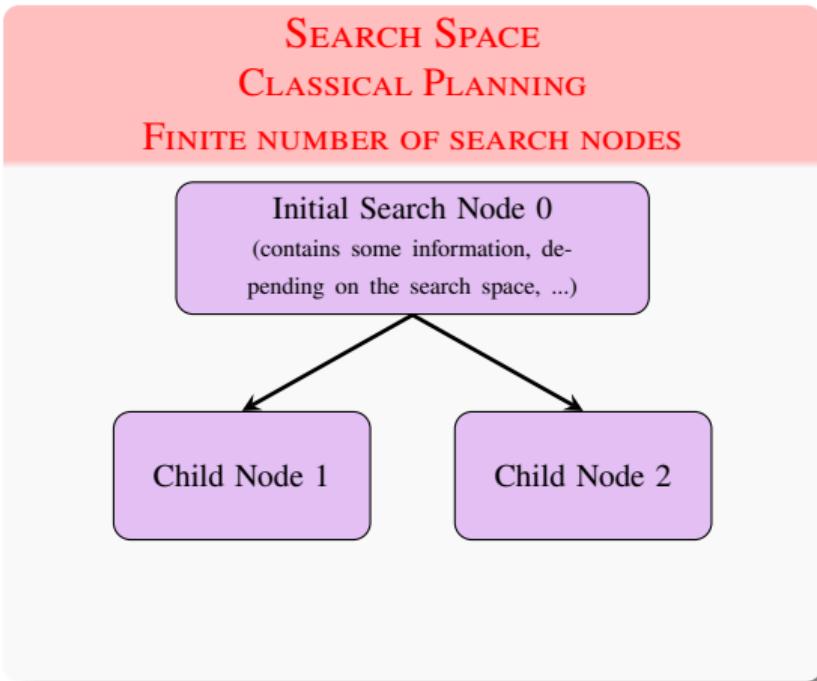


3) A SUCCESSOR FUNCTION/BRANCHING RULE

Returning **all successors** (neighbors) of any search node

- **Edges** could correspond to actions (state space search) or to something different (POCL)!
- **Expand** a node corresponds to generate **all** its successors

PLANNING AS SEARCH (CONT.)



- To know when we succeeded

4) A SOLUTION CRITERION

To detect when a node corresponds to a solution

5) A PLAN EXTRACTION

To tell/extract which **plan** a solution node corresponds to

- To decide **how to search**

6) A SEARCH STRATEGY

To choose which node to expand next

PLANNING AS SEARCH (CONT.)

- General Search-Based Planning Algorithm

```
function SEARCH(problem)
    initial-node ← MAKE-INITIAL-NODE(problem)
    open ← {initial-node}
    while (open ≠ ∅) do
        node ← SEARCH-STRATEGY-REMOVE-FROM(open)
        if IS-SOLUTION(node) then
            return EXTRACT-PLAN-FROM(node)
        end if
        for each newnode ∈ SUCCESSORS(node) do
            open ← open ∪ {newnode}
        end for
    end while
    return Failure
end function
```

→ [2]

→ [6]

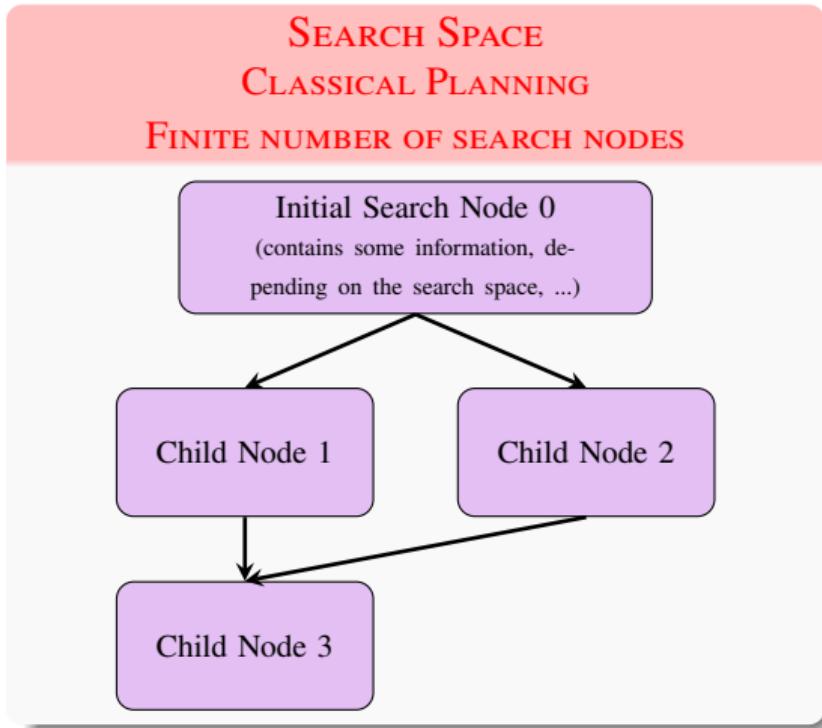
→ [4]

→ [5]

→ [3]

→ *Expanded the entire search space without finding a solution*

SEARCHING GRAPH

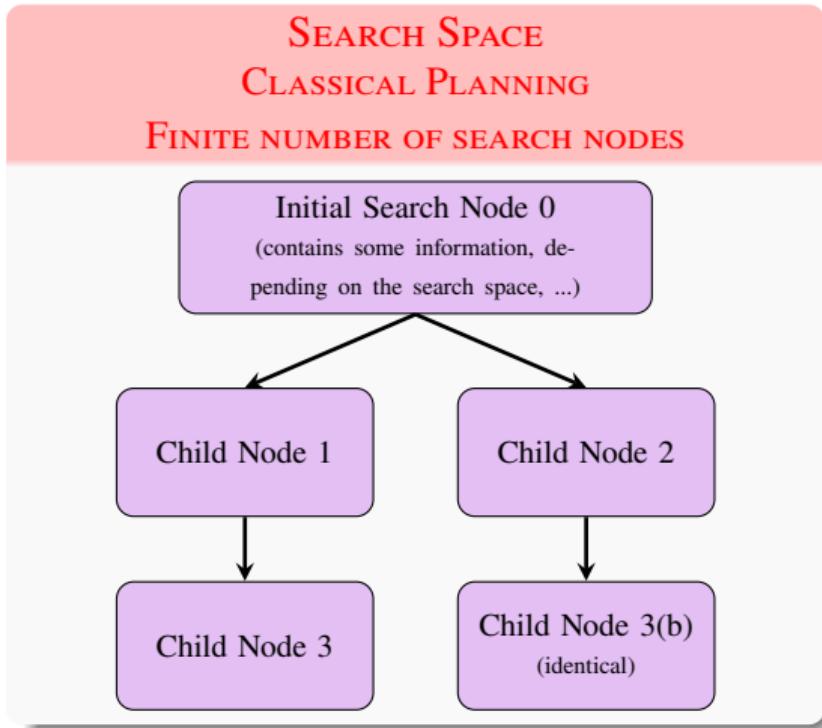


In a graph two nodes can share a successor!

OPTION 1

- Keep track of all visited nodes
- Detect when the same successor is generated again
 - Requires a lot of memory
 - Only investigate a given node once
 - Second time do not expand it!

SEARCHING GRAPH (CONT.)



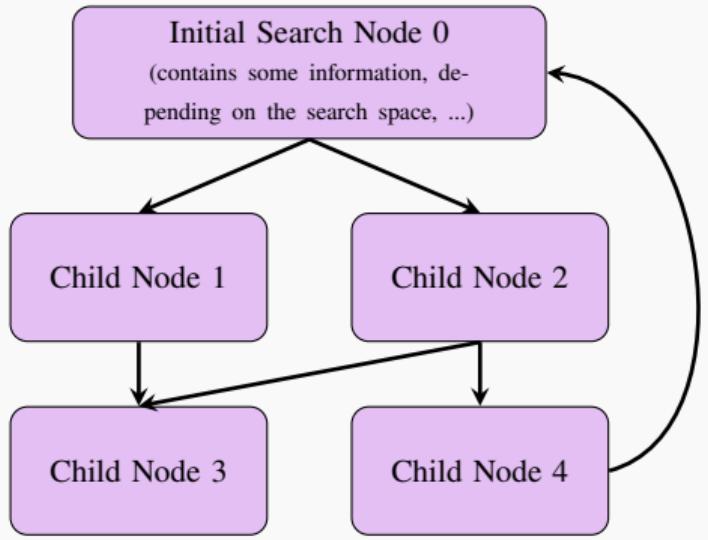
OPTION 2

- Do not keep track of visited nodes!
 - Saves memory
 - Investigate some subtrees multiple times
 - Search space is visited as a *tree*

SEARCHING GRAPH (CONT.)

SEARCH SPACE CLASSICAL PLANNING

FINITE NUMBER OF SEARCH NODES



LOOPS IN THE SEARCH

- An ancestor may also be a successor
 - \Rightarrow loops in the search graph
- Depending on the search **strategy** it may or may not be necessary to detect and handle this!

SEARCHING GRAPH (CONT.)

- To avoid searching sub-graphs twice (shared successors + loops):

```

function SEARCH(problem)
    initial-node ← MAKE-INITIAL-NODE(problem) → [2]
    open ← {initial-node}
    added ← {initial-node}
    while (open ≠ ∅) do
        node ← SEARCH-STRATEGY-REMOVE-FROM(open)
        if IS-SOLUTION(node) then
            return EXTRACT-PLAN-FROM(node) → [6] // Nodes removed from open, but not from added!
        end if
        for each newnode ∈ SUCCESSORS(node) do
            if newnode ∉ added then → [4]
                open ← open ∪ {newnode}
                added ← added ∪ {newnode} → [5]
            end if
        end for
    end while
    return Failure → [3]
end function

```

→ Expanded the entire search space without finding a solution

ASPECTS OF SEARCH

DEFINED BY THE SEARCH SPACE

- 1) Node structure
- 2) Generating initial search node
- 3) Branching rule, creating successors
- 4) Determining if a node is a solution
- 5) Extracting a plan from a node

Forward State Space, Backward Goal Space,
Partial Order Causal Link, ...

DEFINED BY THE SEARCH STRATEGY

Uninformed

- Depth first (DFS)
- Breadth first (BFS)
- Dijkstra
- Uniform cost
- Depth limited DFS
- Iter. Deep. DFS
- ...

Informed

- Greedy Best First
- A*
- Weighted A*
- Iter. Deep. A*
- Beam Search
- Hill Climbing (HC)
- Enforce HC
- Simul. Annealing
- ...



Heuristics!

Independence!

PLANNING AS SEARCH: SUMMARY

Search Spaces

- Forward State Space
- Backward Goal Space
- Partial Order Causal Link
- Hierarchical Task Networks

Search Strategies

- Hill Climbing
- Enforced Hill Climbing
- A*
- (Repeated) Weighted A*

Heuristics

- Goal count
- Landmarks
- Pattern Databases
- Relaxation, Delete Relaxation
- Relaxed Planning Graphs

Tweaking Search Space

- Predicates vs State variables
- Lifted Search Space

Tweaking Search Strategies

- Helpful Actions / Preferred Operators
- Dual Queues, Boosted Dual Queues
- Lazy Search

Meta Search Strategies

- Portfolio Planning

PREPARATION FOR HANDS-ON SESSION NEXT WEEK

- We will create and solve planning problems
- On your laptops!
 - Experiment with command line planners
 - Planutils (several pre-compiled planners on Linux)
<https://pypi.org/project/planutils/>
<https://github.com/AI-Planning/planutils>
 - Fast Downward¹ <https://www.fast-downward.org/>
 - Experiment with online planners
 - <http://editor.planning.domains>

¹Instructions to compile it for Linux/Window/MacOS X <https://github.com/aibasel/downward/blob/main/BUILD.md>

SOME TIPS FOR PLANUTILS

- To install pip or pip3 (in Debian like Linux distributions e.g. Ubuntu):
 - Run `apt-get install python-pip` or `apt-get install python3-pip`
- Run `planutils list` for a list of available planners
- To install a planner run `<name>` using names from previous command
- You need python version $\geq 3.6.*$ in order to successfully install and then run `planutils`.
- Planutils requires the `singularity` package to be installed on the machine (a kind of lightweight Docker container).
- If the `pip3 install planutils` does not install also `singularity`, you can install it following the instructions you can find at: https://docs.sylabs.io/guides/latest/user-guide/quick_start.html#quick-installation-steps
- Singularity runs on Linux natively and can also be run on Windows and Mac through virtual machines (VMs) or now natively.

SOME TIPS FOR PLANUTILS (CONT.)

- At this URL: <https://docs.sylabs.io/guides/latest/admin-guide/installation.html> you can find instructions to install singularity on a Linux machine.
- At this URL are also reported instructions to install singularity on Windows and Mac Os X
 - Notice that, in Windows and MacOX you need to install then planutils from the virtual machine that is installed with the instructions at the link.
- If for some reason you have PYTHONHOME or PYTHONPATH set, this may interfere with the singularity images (that are not completely sandboxed). To avoid this, then you can edit the files in ~/.planutils/packages/<planner>/run changing planner with one of the installed planners (e.g. downward) adding option -e to run i.e.
`singularity run -e $(dirname $0)/downward.sif`
- Under Linux singularity do not mount all the host files by default.
 - Edit file /usr/local/etc/singularity/singularity.conf (or /etc/singularity/singularity.conf) changing mount hostfs = no to mount hostfs = yes.

REFERENCES I

- [1] Hector Geffner and Blai Bonet. *A Concise Introduction to Models and Methods for Automated Planning*. Synthesis Lectures on Artificial Intelligence and Machine Learning. Morgan & Claypool Publishers, 2013. ISBN 9781608459698. doi: 10.2200/S00513ED1V01Y201306AIM022. URL <https://doi.org/10.2200/S00513ED1V01Y201306AIM022>.
- [2] Malik Ghallab, Dana S. Nau, and Paolo Traverso. *Automated planning - theory and practice*. Elsevier, 2004. ISBN 978-1-55860-856-6.
- [3] Malik Ghallab, Dana S. Nau, and Paolo Traverso. *Automated Planning and Acting*. Cambridge University Press, 2016. ISBN 978-1-107-03727-4. URL <http://www.cambridge.org/de/academic/subjects/computer-science/artificial-intelligence-and-natural-language-processing/automated-planning-and-acting?format=HB>.
- [4] Patrik Haslum, Nir Lipovetzky, Daniele Magazzeni, and Christian Muise. *An Introduction to the Planning Domain Definition Language*. Synthesis Lectures on Artificial Intelligence and Machine Learning. Morgan & Claypool Publishers, 2019. doi: 10.2200/S00900ED2V01Y201902AIM042. URL <https://doi.org/10.2200/S00900ED2V01Y201902AIM042>.