

COURSE "AUTOMATED PLANNING: THEORY AND PRACTICE"

CHAPTER 17: PLANNING IN NON DETERMINISTIC DOMAINS VIA MODEL CHECKING

Teacher: **Marco Roveri** - marco.roveri@unitn.it
M.S. Course: Artificial Intelligence Systems (LM)
A.A.: 2025-2026
Where: DISI, University of Trento
URL: <https://shorturl.at/A81hf>



Last updated: Wednesday 12th November, 2025

TERMS OF USE AND COPYRIGHT

USE

This material (including video recording) is intended solely for students of the University of Trento registered to the relevant course for the Academic Year 2025-2026.

SELF-STORAGE

Self-storage is permitted only for the students involved in the relevant courses of the University of Trento and only as long as they are registered students. Upon the completion of the studies or their abandonment, the material has to be deleted from all storage systems of the student.

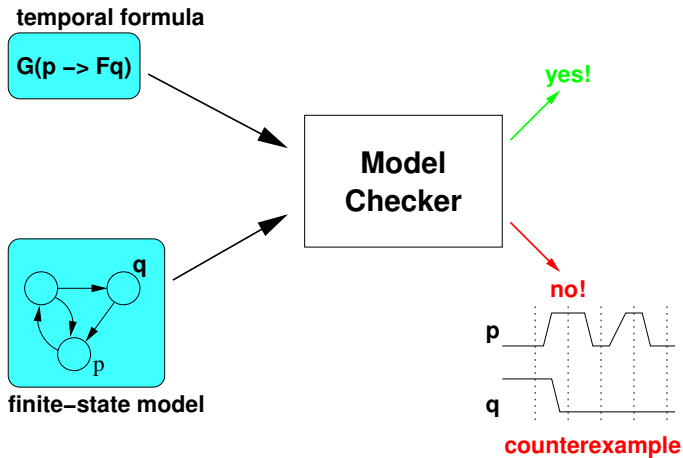
COPYRIGHT

The copyright of all the material is held by the authors. Copying, editing, translation, storage, processing or forwarding of content in databases or other electronic media and systems without written consent of the copyright holders is forbidden. The selling of (parts) of this material is forbidden. Presentation of the material to students not involved in the course is forbidden. The unauthorised reproduction or distribution of individual content or the entire material is not permitted and is punishable by law.

The material (text, figures) in these slides is authored by Marco Roveri.

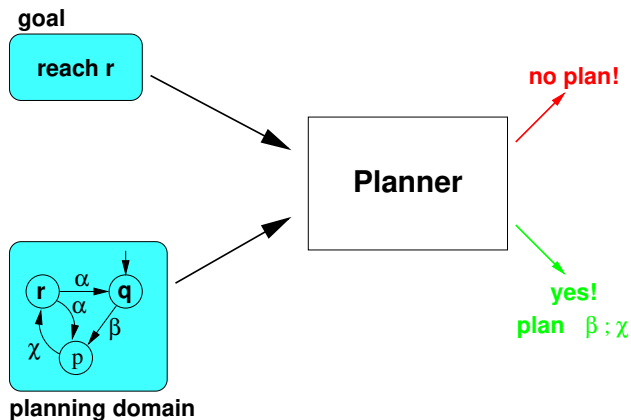
MODEL CHECKING

Model Checking: a technique to *validate a formal model* of a system against a *logical specification*.



PLANNING AS MODEL CHECKING

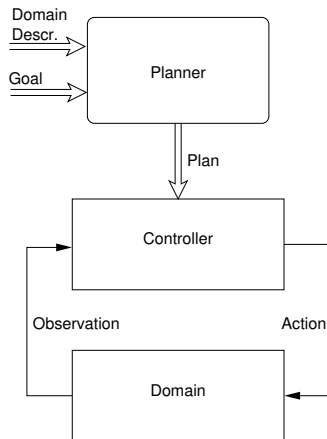
Planning as Model Checking: a technique to *synthesize* a plan from a *formal model of a domain* and a *logical specification of a goal*.



PLANNING

- “Planning” is the problem of building a suitable *plan* for achieving a given *goal* on a given dynamic domain.
- The plan describes the actions to be executed in the domain in order to achieve the goal.
- Different approaches to planning differ for:
 - the assumptions on the *domain*.
 - the assumptions on the *goals*.
 - the assumptions on the *plans*.
 - the techniques used in the *planning algorithm*.

THE CONCEPTUAL MODEL



- The *domain* evolves according to performed *actions*.
- The *controller* provides the *actions* according to the *observations* on the domain and to a *plan*.
- The *planner* synthesizes a *plan* from a *domain description* and a *goal*.

CLASSICAL PLANNING: ASSUMPTIONS

Classical planning relies on simplifying assumptions:

- *finite*: the domain has a finite set of different states.
- *implicit time*: actions are instantaneous state transitions.
- *deterministic*: the initial state is completely specified; each action, if applicable in a state, brings to a single new state.
- *observable*: the state of the system is fully observable to the controller.
- *basic goals*: goals are sets of target states; the objective is to build a plan that leads to one of the goal states.
- *off-line planning*: the plan is built once and for all, before its execution.
- ...

CLASSICAL PLANNING

- A plan is a *sequence of actions* whose execution leads from the initial state of the domain to a goal state.
 - the evolution is completely determined by the executed actions (deterministic domain).
 - no feedback is needed from the domain (observation is not necessary).
- Main difficulty:
 - in spite of simplifying assumptions, search space is huge for practical domains.

NON-DETERMINISM

The restriction to *deterministic* domains is unrealistic in many practical cases:

- There may be uncertainty on the *initial state*.
- *Actions* may have *non-nominal outcomes* that are highly critical.
- There may be *exogenous events* or *non-controllable actions*.

In *non-deterministic domains*:

- More than one initial state is possible.
- An action executed in a state may lead to many different states.

Main difficulties:

- Plans cannot simply be sequences of actions.
- Domain models are more complex.
- Lifting most classical planning techniques does not work in practice.

NON-DETERMINISM: PROBABILISTIC OUTCOMES

It is often possible to associate probabilities to the outcomes of actions:

- non-nominal outcomes of actions often have a very small probability.
- when throwing a fair dice, all the outcomes have equal probability.
- exogenous events occurs with a given probability.

In planning with probabilistic domains (e.g., MDP-based planning):

- probabilities are associated to the outcomes of actions.
- goals are represented by utility, or value functions.
- planning consists in searching for a plan maximizing the utility function.

Main difficulty:

- *optimality* comes into play.

We will not discuss probabilistic models!

PARTIAL OBSERVABILITY

In several realistic problems the domain status is only partially visible at run-time to the controller:

- different states of the domain are *indistinguishable* for the controller.
- certain information on the domain is available only after some *sensing actions*.

Particular cases:

- *full observability*.
- *null observability* (conformant planning).

Main difficulty:

- for non-deterministic domains, the controller has a partial knowledge of the domain status.

EXTENDED GOALS

The assumption that planning goals are sets of final desired states is unrealistic in many cases:

- *safety conditions* (states to be avoided) may complement the main goal.
- *reactive systems* (infinite plans that react to changes).
- *preferences* among plans and *search control rules* can be specified.

Temporally extended goals express conditions on the whole, possibly infinite sequences of states resulting from plan execution.

Main difficulty:


- the plan should trace past history (different “active” goals depending on past events).

PLANS


Different kinds of plans are necessary for dealing with different assumptions on domains and goals.

- *sequences of actions* are sufficient for *classical planning*.
- *conditional plans* are necessary for dealing with *non-deterministic* domains under *partial observability*.
- *history-dependent plans* are necessary for *extended goals*.
- *explicit time representation* may be necessary in plans.
- ...

PLANNING AS MODEL CHECKING



IN THIS LECTURE



REMARKS ON PLANNING SYSTEM...

...must not be limited to plan synthesis.

- Plan validation: verification that a *plan* satisfies a given *property* (not necessarily a goal...).
- Plan simulation: interactive simulation of the executions of a *plan* on a given *domain*.
- Plan synthesis: automatic generation of the plan from a description of the *domain* and from the *goal*.
- ...

MODEL CHECKING

- Model Checking is a widely used automatic technique for verifying design and for discovering possible bugs.

In this lecture . . .

- A mild introduction to model checking.
- Symbolic representation of Finite State Machine (FSM) based on Binary Decision Diagrams (BDDs).
- An introduction to symbolic model checking algorithms for Computation Tree Logic (CTL).

MODEL CHECKING (CONT.)

- Model Checking is a verification technique that solves the problem

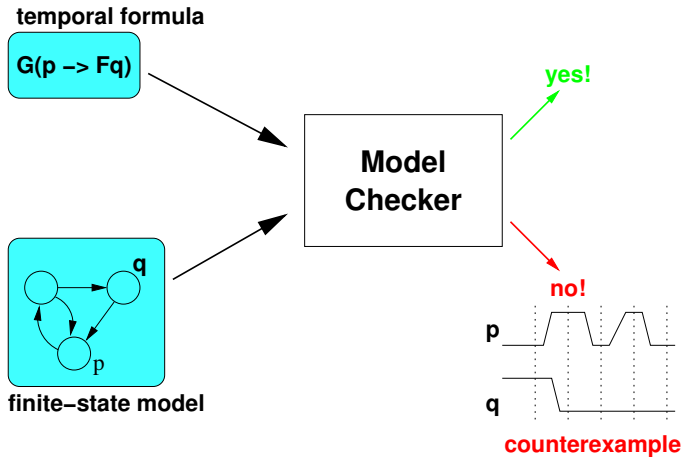
$$\mathcal{M} \models \varphi$$

where. . .

- Model \mathcal{M} is represented as a FSM.
- Property φ is a temporal logic formula.
- Model checking algorithms exhaustively traverse the model guided by the property.

MODEL CHECKER


- Software tool for system verification.



FORMAL MODEL: FSM



- A model is defined as a tuple $\mathcal{M} = \langle S, S_0, T, P, L \rangle$, where:

- S is a finite set of states.
- S_0 is the initial set of states.
- P is a finite set of atomic propositions.
- $T \subseteq S \times S$ is the transition relation.
- $L : P \mapsto 2^S$ assigns to each proposition $p \in P$ the set of states where p holds.




COMPUTATION TREES

- Computation model is a *tree*:




TEMPORAL LOGIC: LTL


- Express properties of all computation paths.
 - Invariantly p : $G p$



- Finally p : $F p$




- p holds until q holds: $p U q$



TEMPORAL LOGIC: CTL


Express properties on computation paths.

- There exists a path where p finally holds: $\mathbf{EF} p$




TEMPORAL LOGIC: CTL (II)

- Property p finally holds for all paths: $\mathbf{AF} p$



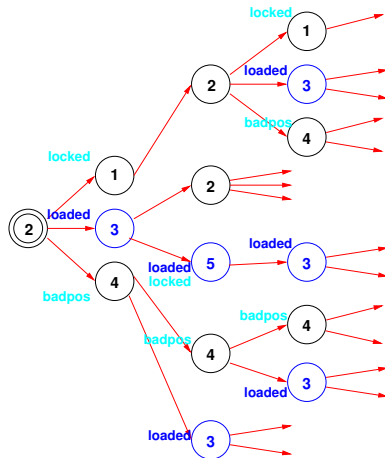

TEMPORAL LOGIC: CTL (III)

- Invariant p holds along all paths: $AG\ p$



CTL EXAMPLES: AG EF

- It is always possible to reach a state where *loaded* holds: **AG EF loaded**



SYMBOLIC MODEL CHECKING

- Symbolic model checking:

- Face the *state explosion problem* of explicit state model checking.
- Use the formula φ to represent the set of states where φ holds.

$$\varphi \rightsquigarrow \{s \mid M, s \models \varphi\}$$

- Use of Boolean formulas to represent sets and relations.
- Fix-point characterization of CTL operators.

$$\mathbf{EF}p \leftrightarrow p \vee \mathbf{EXEF}p$$

$$\mathbf{E}[p\mathbf{U}q] \leftrightarrow (q \vee (p \wedge \mathbf{EX}(\mathbf{E}[p\mathbf{U}q])))$$

$$\mathbf{AG}p \leftrightarrow p \wedge \mathbf{AXAG}p$$

$$\mathbf{EG}p \leftrightarrow p \wedge \mathbf{EXEG}p$$

$$\mathbf{EF}p = \text{lfp}_Z\{p \cup \mathbf{EX}Z\}$$

$$\mathbf{E}[p\mathbf{U}q] = \text{lfp}_Z\{q \cup (p \cap \mathbf{EX}Z)\}$$


$$\mathbf{AG}p = \text{gfp}_Z\{p \cap \mathbf{AX}Z\}$$

$$\mathbf{EG}p = \text{gfp}_Z\{p \cap \mathbf{EX}Z\}$$

... ..


FIX-POINT CHARACTERIZATION OF CTL OPERATORS

- **EF** *loaded*



FIX-POINT CHARACTERIZATION OF CTL OPERATORS (II)

- **EG** *loaded*



SYMBOLIC REPRESENTATION

- A vector \mathbf{x} of Boolean variables where each Boolean variable corresponds to an atomic proposition in P .

$$\mathbf{x} = \{locked, badpos, loaded\}$$

- A state s is represented with a formula $\xi(s)$ on the propositions:

$$\textcircled{1} \rightsquigarrow locked \wedge \neg badpos \wedge \neg loaded$$

$$\textcircled{3} \rightsquigarrow \neg locked \wedge \neg badpos \wedge loaded$$

$$\textcircled{5} \rightsquigarrow locked \wedge \neg badpos \wedge loaded$$

- A set of states $Q \subseteq S$ represented symbolically as:

$$\xi(Q) = \bigvee_{s \in Q} \xi(s)$$

SYMBOLIC REPRESENTATION (II)

- A formula φ represents the set of states where the formula holds:

<i>loaded</i>	$\leadsto \{\textcircled{3}, \textcircled{5}\}$
$\neg \textit{badpos}$	$\leadsto \{\textcircled{1}, \textcircled{2}, \textcircled{3}, \textcircled{5}\}$
<i>loaded</i> \wedge <i>locked</i>	$\leadsto \{\textcircled{5}\}$
<i>badpos</i> \wedge <i>locked</i> \wedge <i>loaded</i>	$\leadsto \emptyset$
<i>loaded</i> \vee <i>locked</i>	$\leadsto \{\textcircled{1}, \textcircled{3}, \textcircled{5}\}$

- Set theoretic operations map to propositional operations:

$$\begin{aligned} \xi(Q_1 \setminus Q_2) &= \xi(Q_1) \wedge \neg \xi(Q_2) & \xi(Q_1 \cup Q_2) &= \xi(Q_1) \vee \xi(Q_2) \\ \xi(Q_1 \cap Q_2) &= \xi(Q_1) \wedge \xi(Q_2) \end{aligned}$$

SYMBOLIC REPRESENTATION (III)

- A new vector \mathbf{x}' of Boolean variables to encode next state:

$$\mathbf{x}' = \{loaded', locked', badpos'\}$$

- A transition $t = \langle s_s, s_d \rangle$ encoded as

$$\xi(t) = \xi(\langle s_s, s_d \rangle) = \xi(s_s) \wedge \xi'(s_d)$$

$$\begin{aligned} \langle \textcircled{1}, \textcircled{2} \rangle &\rightsquigarrow \left(\begin{array}{c} (locked \wedge \neg badpos \wedge \neg loaded) \\ (\neg locked' \wedge \neg badpos' \wedge \neg loaded') \end{array} \wedge \right) \\ \langle \textcircled{2}, \textcircled{3} \rangle &\rightsquigarrow \left(\begin{array}{c} (locked \wedge \neg badpos \wedge \neg loaded) \\ (\neg locked' \wedge \neg badpos' \wedge loaded') \end{array} \wedge \right) \end{aligned}$$

- Transition relation T represented symbolically as:

$$\xi(T) = \bigvee_{t \in T} \xi(t)$$

SYMBOLIC EXPLORATION OF THE MODEL

- Images work on sets of states:

- The forward image $Flmg(S)$ of a set S is

$$Flmg(S) = \{s' \mid s \in S \wedge \langle s, s' \rangle \in T\} \quad \leadsto \quad \exists \mathbf{x}. (S(\mathbf{x}) \wedge T(\mathbf{x}, \mathbf{x}'))$$

- The backward image $Blmg(S)$ of a set S is

$$Blmg(S) = \{s \mid s' \in S \wedge \langle s, s' \rangle \in T\} \quad \leadsto \quad \exists \mathbf{x}'. (S(\mathbf{x}') \wedge T(\mathbf{x}, \mathbf{x}'))$$

SYMBOLIC CTL MODEL CHECKING

- The set of states where a formula φ holds is represented symbolically.
- $M, S_0 \models \varphi$ reduces to $(\neg\varphi \wedge S_0) = \perp$
- Basic CTL operations represented as

$$\mathbf{EX}(\varphi) = \exists \mathbf{x}'. (T(\mathbf{x}, \mathbf{x}') \wedge \varphi(\mathbf{x}'))$$

$$\mathbf{AX}(\varphi) = \forall \mathbf{x}'. (T(\mathbf{x}, \mathbf{x}') \rightarrow \varphi(\mathbf{x}'))$$

- Fix-point computations as propositional transformations.

$$\mathbf{AF} \varphi = \text{lfp}_Z \{ \varphi \cup \mathbf{AX}Z \}$$

$$S_0 = \perp$$

$$S_1 = \varphi \vee \mathbf{AX}S_0 = \varphi$$

$$S_2 = \varphi \vee \mathbf{AX}S_1 = \varphi \vee \mathbf{AX}\varphi$$

$$S_3 = \varphi \vee \mathbf{AX}S_2 = \varphi \vee \mathbf{AX}(\varphi \vee \mathbf{AX}\varphi)$$




...

$$S_i = S_{i-1}$$

BINARY DECISION DIAGRAMS

- The problem in symbolic model checking:
 - The need for efficient and practical representation and manipulation of propositional formulae.
- **Binary Decision Diagrams** (BDDs).
 - Canonical form for propositional formulae.
 - Efficiently managed by BDD packages.



BINARY DECISION DIAGRAMS (II)

 $\{\textcircled{3}\}$  $\{\textcircled{1}, \textcircled{3}, \textcircled{5}\}$  $\{\langle \textcircled{4}, \textcircled{4} \rangle, \langle \textcircled{4}, \textcircled{3} \rangle\}$ 

BINARY DECISION DIAGRAMS (III)

- BDDs variable ordering is important.

$$(a_1 \leftrightarrow b_1) \wedge (a_2 \leftrightarrow b_2) \wedge (a_3 \leftrightarrow b_3)$$



BINARY DECISION DIAGRAMS: PROS AND CONS

● Pros

- Equality test of two BDDs in constant time.
- Basic propositional operations polynomial on the size of operands.

$$O(\varphi_1 \otimes \varphi_2) = O(|\varphi_1| |\varphi_2|)$$

● Cons

- Quantification exponential in the variables being quantified.
- BDD size may be exponential in the number of variables.
 - e.g., multiplier.

- They work well in several practical applications.

INTRODUCTION

The “*Planning as Model Checking*” approach defines a general framework for dealing with:

- non-deterministic domains...
- extended goals...
- partial observability...
- ...and their combination

In “*Planning as Symbolic Model Checking*” BDD-based symbolic techniques are used for:

- a compact representation of domains and plans
- an efficient search in the domain in the planning algorithm


LECTURE OVERVIEW

In this part of the lecture:

- ➊ General framework
- ➋ Reachability goals
- ➌ Extended goals
- ➍ Partial observability

PLANNING DOMAIN: FORMAL DEFINITIONS

- A *planning domain* is defined as a tuple $\mathcal{D} = \langle S, S_0, A, T, P, L, O, X \rangle$:
 - S is a finite set of states.
 - S_0 is the initial set of states.
 - A is a finite set of actions.
 - $T \subseteq S \times A \times S$ is the transition relation.
 - P is a finite set of atomic propositions.
 - $L : P \mapsto 2^S$ assigns to each $p \in P$ the set of states where p holds.
 - O is a finite set of possible observations.
 - $X \subseteq S \times O$ is the observation relation.



PLANNING DOMAINS: SYMBOLIC REPRESENTATION

- States represented as in symbolic model checking: $\xi(\mathbf{s})$.
- A vector α of Boolean variables, each corresponding to an action in A :

$$\alpha = \{GoWest, GoEast, GoNorth, GoSouth\}$$

- An action represented by assigning *True* to the corresponding Boolean variable:

$$\xi(GoWest) \rightsquigarrow GoWest \quad \xi(GoSouth) \rightsquigarrow GoSouth$$

- Alternatively, compact logarithmic encoding of actions.
- A transition $t = \langle \mathbf{s}_s, \mathbf{a}, \mathbf{s}_d \rangle$ encoded as

$$\xi(t) = \xi(\langle \mathbf{s}_s, \mathbf{a}, \mathbf{s}_d \rangle) = \xi(\mathbf{s}_s) \wedge \xi(\mathbf{a}) \wedge \xi'(\mathbf{s}_d)$$

PLANNING DOMAINS: SYMBOLIC REPRESENTATION (II)

- Observations encoded with a vector \mathbf{o} of Boolean variables.

$$\mathbf{o} = \{O_{WallN}, O_{WallS}, O_{WallE}, O_{WallO}\}$$

- An observation represented by assigning *True* to the corresponding Boolean variable.

$$\xi(O_{WallW}) \rightsquigarrow O_{WallW} \qquad \xi(O_{WallS}) \rightsquigarrow O_{WallS}$$

SYMBOLIC EXPLORATION OF THE DOMAINS

- Images work on set of states:

- Forward: set of states reachable from a set of states S


$$FImg(S) = \{s' \mid s \in S \wedge \langle s, a, s' \rangle \in T\} \rightsquigarrow \exists \mathbf{x}.\alpha.(S(\mathbf{x}) \wedge T(\mathbf{x}, \alpha, \mathbf{x}'))$$

- Backward: set of states from which a set of states S is reachable

$$BImg(S) = \{s \mid s' \in S \wedge \langle s, a, s' \rangle \in T\} \rightsquigarrow \exists \mathbf{x}'.\alpha.(S(\mathbf{x}') \wedge T(\mathbf{x}, \alpha, \mathbf{x}'))$$

PLANS: FORMAL MODEL

- A *plan* Π for a planning domain \mathcal{D} is a tuple $\Pi = \langle \Sigma, \sigma_0, \Psi, \Upsilon \rangle$ where:
 - Σ is a finite set of *plan states*.
 - σ_0 is the *initial* plan state.
 - $\Psi : \Sigma \times \mathcal{O} \mapsto \mathcal{A}$ associates to a pair $\langle \sigma, o \rangle$ an action α to execute.
 - $\Upsilon : \Sigma \times \mathcal{O} \mapsto \Sigma$ associates to a pair $\langle \sigma, o \rangle$ a new plan state σ' .



- A configuration for a domain \mathcal{D} and a plan Π is a pair $\langle \mathbf{s}, \sigma \rangle \in \mathcal{S} \times \Sigma$.
- Plan execution is represented with configuration transitions:

$$\langle \mathbf{s}, \sigma \rangle \rightarrow \langle \mathbf{s}', \sigma' \rangle$$

MOTIVATIONS

Reachability goals:


- First and “basic” planning problem:
 - find a plan to achieve a desired final situation.
- When tackled within classical planning:
 - solutions are sequences of actions.
 - solutions are guaranteed to achieve the goal.

Adding non-determinism:

- Plans as sequences of actions are bound to failure.
 - Initial situation might be uncertain.
 - Actions can have non-deterministic effects.
- Plans of different strength, e.g. weak, strong.
- Plans of different structure: sense current state and execute an action.

EXAMPLE

- The problem ...
 - ... find a plan from a (set of) initial state(s) to a goal set of states.



AIM

- Objectives:

- Allow for planning under *full observability* in *non-deterministic* domains.
- *Synthesis* of plans of different *strength*.
- Deal *in practice* with *large size domains*.

- Problems:


- How can we handle non-deterministic actions?
- Which kind of plans must be generated?
- Which planning algorithms?
- How can planning algorithms handle large domains?

THE PMC APPROACH FOR REACHABILITY GOALS

- **Solution plans** encode **conditional behavior** based on domain state.
- **Planning algorithms** exploit . . .
 - symbolic representation of the domain.
 - symbolic representation of solution plans.


EXAMPLE

- **Strong Solutions:** plans that are guaranteed to reach the goal.
 - All executions traces are acyclic and reach the goal.




EXAMPLE

- **Weak Solutions:** plans that may achieve the goal.
 - At least one execution trace reaches the goal.



EXAMPLE

- **Strong Cyclic Solutions:** trial and error strategies.
 - The goal is reachable from all states of the execution traces.
 - Solutions are guaranteed to reach the goal under the assumption of “fair” execution.



STATE-ACTION TABLES

- Solutions are **memory-less** plans that map states to actions to execute.
- Such solutions can be represented as “**state-action tables**” (SA).

State	Action
R1	GoUp
R2	GoRight
R3	GoDown
R4	GoRight

- State-action tables map in a general plan. . .
 - . . . without plan states.
 - . . . switching on the current state.

SYMBOLIC REPRESENTATION OF SA TABLES

- A state-action pair $p = \langle s, a \rangle$ is represented symbolically as:

$$\xi(p) = \xi(s) \wedge \xi(a)$$


- A state-action table π is encoded symbolically as:

$$\xi(\pi) = \bigvee_{p \in \pi} \xi(p)$$

- Set operations on state-action tables performed as:

$$\begin{aligned}\xi(\pi_1 \setminus \pi_2) &= \xi(\pi_1) \wedge \neg \xi(\pi_2) & \xi(\pi_1 \cup \pi_2) &= \xi(\pi_1) \vee \xi(\pi_2) \\ \xi(\pi_1 \cap \pi_2) &= \xi(\pi_1) \wedge \xi(\pi_2)\end{aligned}$$

STRONGPLAN: ALGORITHM INTUITION



STRONGPLAN: ALGORITHM

```

function STRONGPLAN ( $l$ ,  $G$ );
   $OSA := FAIL$ ;  $SA := \emptyset$ ;
  while ( $OSA \neq SA$ ) do
     $PIMG := SPIMG(G \cup StOf(SA))$ ;
     $NSA := PRUNESTATES(PIMG, G \cup StOf(SA))$ ;
     $OSA := SA$ ;
     $SA := SA \cup NSA$ ;
  done;
  if ( $l \subseteq (G \cup StOf(SA))$ ) then
    return  $SA$ ;
  else
    return  $FAIL$ ;
end;

```

```

function AF ( $l$ ,  $G$ );
   $OS := FAIL$ ;  $S := \emptyset$ ;
  while ( $OldS \neq S$ ) do
     $PIMG := AX(G \cup S)$ ;

     $OS := S$ ;
     $S := S \cup PIMG$ ;
  done;
  if ( $l \subseteq (G \cup S)$ ) then
    return  $OK$ ;
  else
    return  $FAIL$ ;
end;

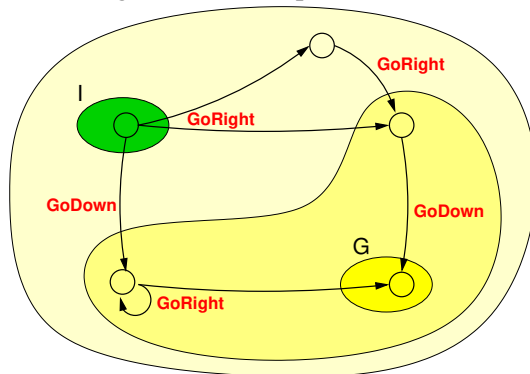
```


STRONGPLAN: PROPERTIES

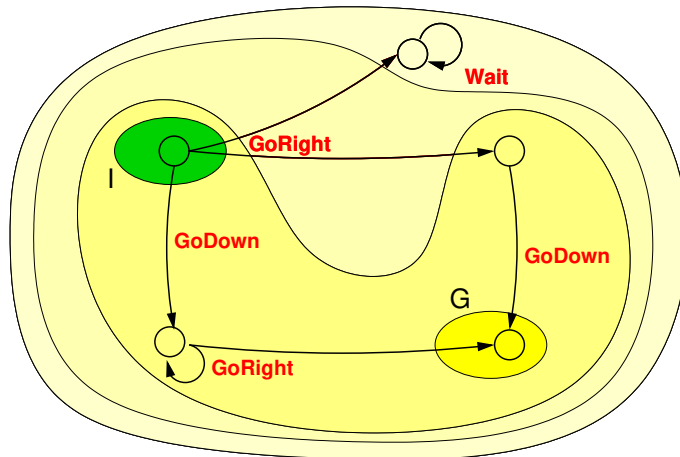
- The algorithm terminates.
- The algorithm is correct and complete.
 - Whenever a strong solution exists it finds a state-action table that is a strong solution.
 - Whenever a strong solution does not exist FAIL is returned.
- The algorithm computes “optimal” solutions.
 - Optimality defined on the length of the plan execution traces.

WEAKPLAN: ALGORITHM INTUITIONS

- WEAKPLAN is similar to the STRONGPLAN algorithm.
 - It performs a *weak pre-image* as basic step.
 - Weak pre-image of a set S computes those state-action pairs whose execution *may* lead to a state $s \in S$.
- WEAKPLAN similar to the SMC algorithm to compute **EF**.



STRONGCYCLICPLAN: ALGORITHM INTUITION



STRONGCYCLICPLAN: ALGORITHM

```
function STRONGCYCLICPLAN ( $l$ ,  $G$ );  
   $OSA := \emptyset$ ;  
   $SA := UNIVSA$ ;  
  while ( $OSA \neq SA$ ) do  
     $OSA := SA$ ;  
     $SA := PRUNEUNCONNECTED(PRUNEOUTGOING(SA, G), G)$ ;  
  done;  
  if ( $l \subseteq (G \cup StOf(SA))$ ) then  
    return REMOVE_NONPROGRESS( $SA, G$ );  
  else  
    return FAIL;  
end;
```

STRONGCYCLICPLAN vs AG EF

- Similarities:
 - Similar meaning (i.e., from all reached states it is possible to achieve the goal).
 - Plan validation performed by verifying **AG EF** p on the synchronous product of domain and plan.
- Differences:
 - The computation of **AG EF** is performed with two *nested distinct* fix point calculations.
 - STRONGCYCLICPLAN *interleaves* the steps of the two fix point calculations.


STRONGCYCLICPLAN: PROPERTIES

- The algorithm terminates.
- The algorithm is correct and complete.
 - Whenever a strong cyclic solution exists it finds a state-action table that is a strong cyclic solution.
 - Whenever a strong cyclic solution does not exist FAIL is returned.

MOTIVATIONS

- Realistic:
 - Environment is seldomly fully observable.
 - Some form of sensing is often available.
- General: full observability, conformancy are special cases.
- Related to relevant problems:
 - Diagnosis.
 - Homing/distinguishing problem.
 - Game theory.

EXAMPLE: A ROBOT-WORLD WITH SENSING



AIM

Objectives:

- Allow for planning under *partial observability*...
- in *non-deterministic domains*.
- Deal *in practice* with domains of *large size*.

Problems:

- How can we express partial observability?
- Which kind of plans must be generated?
- Planning algorithms?
- How can planning algorithms deal with large domains?

THE PMC APPROACH FOR PO


- **Partial observability** expressed by **formulae** that encode observation relation.
- **Plans** encode **conditional behavior** based on sensing.
- **Planning algorithms** exploit **symbolic representation of observations** (as well as actions and states).
- **Experimental results** show that algorithms work **in practice**.

GOALS

- Due to nondeterminism and partial observability, the controller has a partial knowledge of the state.
- We consider *strong reachability* goals:
 - At the end of plan execution, the executor *knows* that the goal has been reached...
 - ...in spite of possibly not knowing exactly the state.


PLANS

- Structure of plans:
 - In general: conditional over observations
 - Conformant case: sequences
 - Full observability case: conditional over states
- Expressed as general FSM:
 - Plan state: “plan program counter”.
 - Branching conditions on:
 - “plan program counter”.
 - observations.
- Validated by model-checking.




PLAN SEARCH: BELIEF STATES

- Due to nondeterminism and partial observability, status known to controller can be uncertain:
 - ...because the initial situation is uncertain...




- ... or because actions may have unpredictable results.




Thus, planning proceeds over sets of states: *belief states (BS)*

PLAN SEARCH: BASE STEPS







- An action transforms a BS into a BS



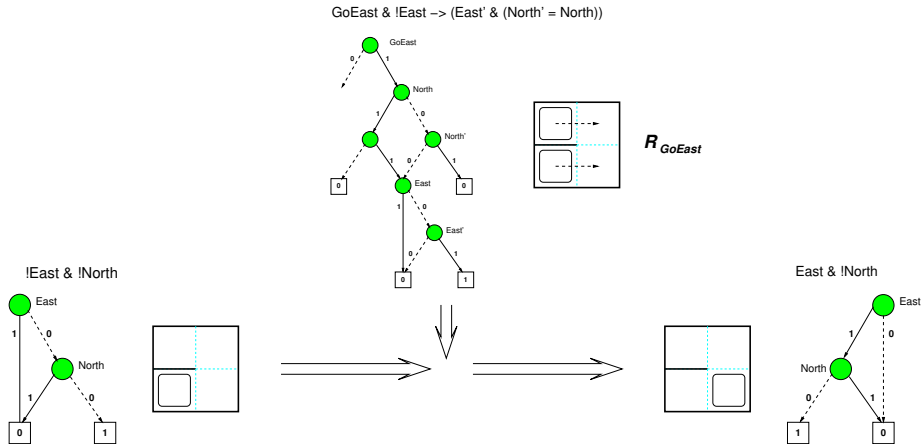
- Observing *may* split a BS



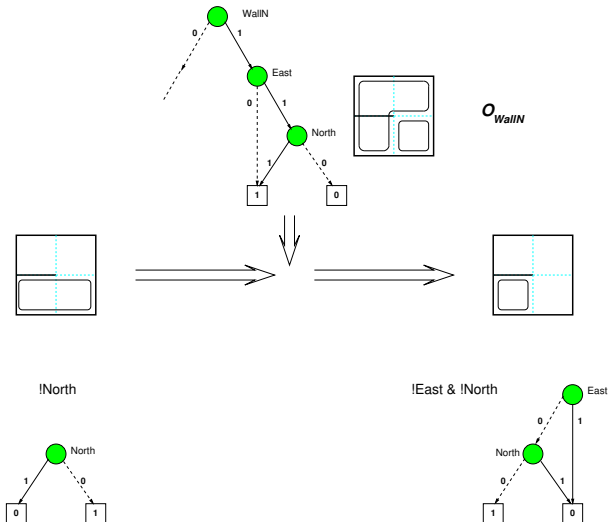
BDD REPRESENTATION OF BS




BDD TRANSFORMATION OF BS: ACTING



BDD TRANSFORMATION OF BS: OBSERVING



SEARCH SPACE




SEARCH ALGORITHM: A DFS SEARCH




SEARCH ALGORITHM: A DFS SEARCH




SEARCH ALGORITHM: A DFS SEARCH



SEARCH ALGORITHM: A DFS SEARCH



SEARCH ALGORITHM: A DFS SEARCH



HEURISTIC SEARCH


Motivations:

- DFS may produce dumb plans.
- DFS may be not efficient.
- Even simple user-provided heuristics may help a lot.
- Heuristics may be extracted from domain/problem.

Example:

- Favouring observations.
- Avoiding “stepping back”.

SEARCH ALGORITHM: A HEURISTIC SEARCH





CONFORMANT ALGORITHMS

Motivations:

- Relevant applications (e.g. reset sequences).
- Ad-hoc highly efficient planning algorithms.
- Dedicated heuristics and fwd/bwd algorithms, e.g.:
 - breadth-first search in belief space
 - knowledge-driven search
 - ...

CONFORMANT ALGORITHMS (II)

- Search space only features “or” branching.
- BFS search in belief space is feasible...
- ...but it's not the only way.



PROPERTIES OF ALGORITHMS

- The algorithms terminate.
- The algorithms are correct and complete:
 - Whenever a strong solution exists they find a plan which is a strong solution.
 - Whenever a strong solution does not exist they return FAIL.

COMBINING EXTENDED GOALS AND PARTIAL OBSERVABILITY

- Realistic domains (robot navigation, embedded controllers) combine *partial observability* and *extended goals*.
- Challenging problem:
 - knowledge goals: introducing “*knows*” in the goal language.
 - plan synthesis requires combining plan contexts and belief states.
 - no trivial combination of existing PO and EG algorithms.
- Work done so far:
 - a language for expressing temporally-extended knowledge goals:
 - example: $\text{AG } p \wedge \text{AF } \text{K } q$
 - a framework for the validation of plans:
 - a monitor is used for epistemic tracing to trace the belief state.


PLANNING IN “ADVERSARIAL” ENVIRONMENT

- So far we have assumed a “fair” environment:
 - all the action outcomes have some chance to occur.
- In several applications:
 - the non-determinism is due to uncontrollable, adversarial agents:
 - we cannot assume that they will execute actions in a fair way.
 - they can apply a strategy to prevent goal achievement.
 - the non-determinism is due to lack of information:
 - some action outcomes may never happen in the “real” domain.
- State of the art:
 - some forms of *adversarial* planning as model checking addressed by Jensen and Veloso [6]
 - partial observability and extended goals are still open
 - non-deterministic effects, partial observability and time are still also open


PLANNING WITH PROBABILITIES

- ADDs (Algebraic Decision Diagrams) can be used for representing:

- expected rewards:



- probabilities of outcomes:



- In SPUDD [5], ADDs are exploited for efficient value iteration in MDP planning.

WRAP-UP

Planning via Model Checking:

- A single, *well founded framework* for a variety of planning problems.
- Several specialized algorithms available for plan synthesis.
- Works in practice.
- Plan validation “for free”: Model Checking of plan against domain.

REFERENCES I

- [1] Hector Geffner and Blai Bonet. *A Concise Introduction to Models and Methods for Automated Planning*. Synthesis Lectures on Artificial Intelligence and Machine Learning. Morgan & Claypool Publishers, 2013. ISBN 9781608459698. doi: 10.2200/S00513ED1V01Y201306AIM022. URL <https://doi.org/10.2200/S00513ED1V01Y201306AIM022>.
- [2] Malik Ghallab, Dana S. Nau, and Paolo Traverso. *Automated planning - theory and practice*. Elsevier, 2004. ISBN 978-1-55860-856-6.
- [3] Malik Ghallab, Dana S. Nau, and Paolo Traverso. *Automated Planning and Acting*. Cambridge University Press, 2016. ISBN 978-1-107-03727-4. URL <http://www.cambridge.org/de/academic/subjects/computer-science/artificial-intelligence-and-natural-language-processing/automated-planning-and-acting?format=HB>.
- [4] Patrik Haslum, Nir Lipovetzky, Daniele Magazzeni, and Christian Muise. *An Introduction to the Planning Domain Definition Language*. Synthesis Lectures on Artificial Intelligence and Machine Learning. Morgan & Claypool Publishers, 2019. doi: 10.2200/S00900ED2V01Y201902AIM042. URL <https://doi.org/10.2200/S00900ED2V01Y201902AIM042>.
- [5] Jesse Hoey, Robert St-Aubin, Alan J. Hu, and Craig Boutilier. SPUDD: stochastic planning using decision diagrams. In Kathryn B. Laskey and Henri Prade, editors, *UAI '99: Proceedings of the Fifteenth Conference on Uncertainty in Artificial Intelligence, Stockholm, Sweden, July 30 - August 1, 1999*, pages 279–288. Morgan Kaufmann, 1999. URL https://dslpitt.org/uai/displayArticleDetails.jsp?mmnu=1&smnu=2&article_id=178&proceeding_id=15. 87
- [6] Rune M. Jensen and Manuela M. Veloso. Obdd-based universal planning for synchronized agents in non-deterministic domains. *J. Artif. Intell. Res.*, 13:189–226, 2000. doi: 10.1613/jair.649. URL <https://doi.org/10.1613/jair.649>. 86